

## Jackson Week 13 IP Part 2

Jackson Kyalo

9/3/2021

### Define the Question

Kira Plastinina is a Russian brand that is sold through a defunct chain of retail stores in Russia, Ukraine, Kazakhstan, Belarus, China, Philippines, and Armenia. The brand's Sales and Marketing team would like to understand their customer's behavior from data that they have collected over the past year. More specifically, they would like to learn the characteristics of customer groups.

### The metric for success

This project will be successful if we are able to determine which individuals are most likely to click on the ads.

### The Outline context

The number of clicks an ad has helps understand how well the ad is being received by its audience. Ads that are targeted to the right audience receive the highest number of clicks. In our case determining the best audience for the ads will help company grow as well as increase the number of clicks and reach.

### Experimental design

1. Define the Questions.
2. Import, load and preview the data.
3. Data Cleaning.
4. Data Analysis.
5. Conclusion and Recommendation.

#### Importing the libraries

```
#Import the data library  
library(data.table)
```

```
## Warning: package 'data.table' was built under R version 4.0.5
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.0.5

## -- Attaching packages ----- tidyverse
1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.3      v dplyr  1.0.7
## v tidyr   1.1.3      v stringr 1.4.0
## v readr   2.0.1      v forcats 0.5.1

## Warning: package 'ggplot2' was built under R version 4.0.5
## Warning: package 'tibble' was built under R version 4.0.5
## Warning: package 'tidyr' was built under R version 4.0.5
## Warning: package 'readr' was built under R version 4.0.5
## Warning: package 'purrr' was built under R version 4.0.5
## Warning: package 'dplyr' was built under R version 4.0.5
## Warning: package 'stringr' was built under R version 4.0.5
## Warning: package 'forcats' was built under R version 4.0.5

## -- Conflicts -----
tidyverse_conflicts() --
## x dplyr::between() masks data.table::between()
## x dplyr::filter()  masks stats::filter()
## x dplyr::first()   masks data.table::first()
## x dplyr::lag()     masks stats::lag()
## x dplyr::last()    masks data.table::last()
## x purrr::transpose() masks data.table::transpose()

library(ggplot2)
library(moments)
```

## Load the dataset

*#Load our data*

```
ecomm=read.csv('http://bit.ly/EcommerceCustomersDataset')
```

## Preview the data

*# preview the head*

```
head(ecomm)
```

```
##   Administrative Administrative_Duration Informational
Informational_Duration
## 1                0                      0          0
0
## 2                0                      0          0
0
## 3                0                      -1          0
```

```

-1
## 4          0          0          0
0
## 5          0          0          0
0
## 6          0          0          0
0
##   ProductRelated ProductRelated_Duration BounceRates ExitRates PageValues
## 1          1          0.000000 0.20000000 0.2000000          0
## 2          2          64.000000 0.00000000 0.1000000          0
## 3          1          -1.000000 0.20000000 0.2000000          0
## 4          2          2.666667 0.05000000 0.1400000          0
## 5         10          627.500000 0.02000000 0.0500000          0
## 6         19          154.216667 0.01578947 0.0245614          0
##   SpecialDay Month OperatingSystems Browser Region TrafficType
## 1          0   Feb                1      1      1          1
## 2          0   Feb                2      2      1          2
## 3          0   Feb                4      1      9          3
## 4          0   Feb                3      2      2          4
## 5          0   Feb                3      3      1          4
## 6          0   Feb                2      2      1          3
##           VisitorType Weekend Revenue
## 1 Returning_Visitor   FALSE   FALSE
## 2 Returning_Visitor   FALSE   FALSE
## 3 Returning_Visitor   FALSE   FALSE
## 4 Returning_Visitor   FALSE   FALSE
## 5 Returning_Visitor    TRUE   FALSE
## 6 Returning_Visitor   FALSE   FALSE

```

### Preview tail

```
tail(ecomm)
```

```

##           Administrative Administrative_Duration Informational
## 12325          0          0          1
## 12326          3         145          0
## 12327          0          0          0
## 12328          0          0          0
## 12329          4          75          0
## 12330          0          0          0
##           Informational_Duration ProductRelated ProductRelated_Duration
BounceRates
## 12325          0          16          503.000
0.000000000
## 12326          0          53          1783.792
0.007142857
## 12327          0          5          465.750
0.000000000
## 12328          0          6          184.250
0.083333333
## 12329          0          15          346.000

```

```

0.000000000
## 12330          0          3          21.250
0.000000000
##          ExitRates PageValues SpecialDay Month OperatingSystems Browser
Region
## 12325 0.03764706    0.00000          0   Nov          2          2
1
## 12326 0.02903061   12.24172          0   Dec          4          6
1
## 12327 0.02133333    0.00000          0   Nov          3          2
1
## 12328 0.08666667    0.00000          0   Nov          3          2
1
## 12329 0.02105263    0.00000          0   Nov          2          2
3
## 12330 0.06666667    0.00000          0   Nov          3          2
1
##          TrafficType      VisitorType Weekend Revenue
## 12325          1 Returning_Visitor   FALSE   FALSE
## 12326          1 Returning_Visitor    TRUE   FALSE
## 12327          8 Returning_Visitor    TRUE   FALSE
## 12328         13 Returning_Visitor    TRUE   FALSE
## 12329         11 Returning_Visitor   FALSE   FALSE
## 12330          2      New_Visitor    TRUE   FALSE

```

### Check the info

```

str(ecomm)

## 'data.frame':    12330 obs. of  18 variables:
## $ Administrative      : int  0 0 0 0 0 0 0 1 0 0 ...
## $ Administrative_Duration: num  0 0 -1 0 0 0 -1 -1 0 0 ...
## $ Informational       : int  0 0 0 0 0 0 0 0 0 0 ...
## $ Informational_Duration: num  0 0 -1 0 0 0 -1 -1 0 0 ...
## $ ProductRelated      : int  1 2 1 2 10 19 1 1 2 3 ...
## $ ProductRelated_Duration: num  0 64 -1 2.67 627.5 ...
## $ BounceRates         : num  0.2 0 0.2 0.05 0.02 ...
## $ ExitRates           : num  0.2 0.1 0.2 0.14 0.05 ...
## $ PageValues          : num  0 0 0 0 0 0 0 0 0 0 ...
## $ SpecialDay          : num  0 0 0 0 0 0 0.4 0 0.8 0.4 ...
## $ Month               : chr  "Feb" "Feb" "Feb" "Feb" ...
## $ OperatingSystems    : int  1 2 4 3 3 2 2 1 2 2 ...
## $ Browser             : int  1 2 1 2 3 2 4 2 2 4 ...
## $ Region              : int  1 1 9 2 1 1 3 1 2 1 ...
## $ TrafficType         : int  1 2 3 4 4 3 3 5 3 2 ...
## $ VisitorType         : chr  "Returning_Visitor" "Returning_Visitor"
"Returning_Visitor" "Returning_Visitor" ...
## $ Weekend             : logi  FALSE FALSE FALSE FALSE TRUE FALSE ...
## $ Revenue             : logi  FALSE FALSE FALSE FALSE FALSE FALSE ...

```

### Check the shape

```
dim(ecomm)

## [1] 12330    18
```

Our code has 1000 rows and 10 columns

## Data Cleaning

### Missing values

```
#check for missing values
sum(is.na(ecomm))

## [1] 112
```

Our data has 112 missing values

```
#check the missing values in each column
colSums(is.na(ecomm))

##      Administrative Administrative_Duration      Informational
##           14              14              14
## Informational_Duration      ProductRelated ProductRelated_Duration
##           14              14              14
##      BounceRates      ExitRates      PageValues
##           14              14              0
##      SpecialDay      Month      OperatingSystems
##           0              0              0
##      Browser      Region      TrafficType
##           0              0              0
##      VisitorType      Weekend      Revenue
##           0              0              0
```

```
#We shall drop the missing values in each columns
df <- na.omit(ecomm)
colSums(is.na(df))

##      Administrative Administrative_Duration      Informational
##           0              0              0
## Informational_Duration      ProductRelated ProductRelated_Duration
##           0              0              0
##      BounceRates      ExitRates      PageValues
##           0              0              0
##      SpecialDay      Month      OperatingSystems
##           0              0              0
##      Browser      Region      TrafficType
##           0              0              0
##      VisitorType      Weekend      Revenue
##           0              0              0
```

## Duplicates

*#Check for duplicates*

```
sum(duplicated(df))
```

```
## [1] 117
```

Our data has 117 duplicated rows. We shall drop all duplicates by selecting only the unique values

*#selecting the unique values*

```
df_new <- unique(df)
```

```
sum(duplicated(df_new))
```

```
## [1] 0
```

*### Identify numeric cols*

```
nums <- unlist(lapply(df_new, is.numeric))
```

```
y <- colnames(df_new[nums])
```

```
y
```

```
## [1] "Administrative" "Administrative_Duration"
```

```
## [3] "Informational" "Informational_Duration"
```

```
## [5] "ProductRelated" "ProductRelated_Duration"
```

```
## [7] "BounceRates" "ExitRates"
```

```
## [9] "PageValues" "SpecialDay"
```

```
## [11] "OperatingSystems" "Browser"
```

```
## [13] "Region" "TrafficType"
```

## Check for outliers

*#Create a dataframe of numeric cols*

```
num <- df_new[y]
```

```
head(num)
```

```
## Administrative Administrative_Duration Informational
Informational_Duration
```

```
## 1 0 0 0
0
```

```
## 2 0 0 0
0
```

```
## 3 0 -1 0
-1
```

```
## 4 0 0 0
0
```

```
## 5 0 0 0
0
```

```
## 6 0 0 0
0
```

```
## ProductRelated ProductRelated_Duration BounceRates ExitRates PageValues
```

```
## 1 1 0.000000 0.20000000 0.2000000 0
```

```
## 2 2 64.000000 0.00000000 0.1000000 0
```


```
## 3 1 -1.000000 0.20000000 0.2000000 0
```

```


## 4          2          2.666667 0.05000000 0.1400000      0
## 5         10         627.500000 0.02000000 0.0500000      0
## 6         19         154.216667 0.01578947 0.0245614      0
##   SpecialDay OperatingSystems Browser Region TrafficType
## 1          0              1      1      1          1
## 2          0              2      2      1          2
## 3          0              4      1      9          3
## 4          0              3      2      2          4
## 5          0              3      3      1          4
## 6          0              2      2      1          3

#Using boxplots to visualize the outliers
for(i in 2:ncol(num)) {
  boxplot(num[i], xlab=colnames(num[i]))
}

```




Administrative\_Duration




Informational







Informational\_Duration




ProductRelated




ProductRelated\_Duration




BounceRates




ExitRates




PageValues




SpecialDay



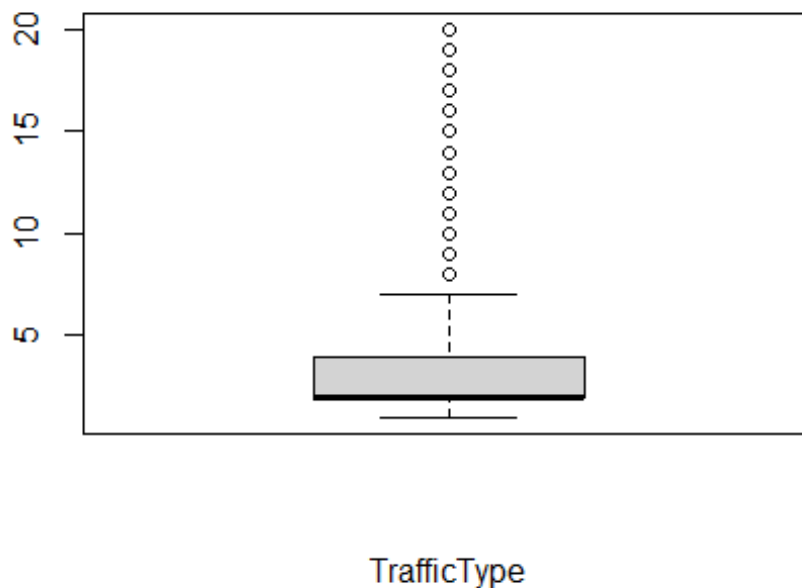
OperatingSystems



Browser



Region



## Data Analysis

### Univariant Analysis

*#Getting the statistical summaries of the data*  
summary(df\_new)

```
## Administrative Administrative_Duration Informational
## Min. : 0.00 Min. : -1.00 Min. : 0.0000
## 1st Qu.: 0.00 1st Qu.: 0.00 1st Qu.: 0.0000
## Median : 1.00 Median : 9.00 Median : 0.0000
## Mean : 2.34 Mean : 81.68 Mean : 0.5088
## 3rd Qu.: 4.00 3rd Qu.: 94.75 3rd Qu.: 0.0000
## Max. : 27.00 Max. : 3398.75 Max. : 24.0000
## Informational_Duration ProductRelated ProductRelated_Duration
## Min. : -1.00 Min. : 0.00 Min. : -1.0
## 1st Qu.: 0.00 1st Qu.: 8.00 1st Qu.: 193.6
## Median : 0.00 Median : 18.00 Median : 609.5
## Mean : 34.84 Mean : 32.06 Mean : 1207.5
## 3rd Qu.: 0.00 3rd Qu.: 38.00 3rd Qu.: 1477.6
## Max. : 2549.38 Max. : 705.00 Max. : 63973.5
## BounceRates ExitRates PageValues SpecialDay
## Min. :0.00000 Min. :0.00000 Min. : 0.000 Min. :0.00000
## 1st Qu.:0.00000 1st Qu.:0.01422 1st Qu.: 0.000 1st Qu.:0.00000
## Median :0.00293 Median :0.02500 Median : 0.000 Median :0.00000
## Mean :0.02045 Mean :0.04150 Mean : 5.952 Mean :0.06197
```

```
## 3rd Qu.:0.01667 3rd Qu.:0.04848 3rd Qu.: 0.000 3rd Qu.:0.00000
## Max. :0.20000 Max. :0.20000 Max. :361.764 Max. :1.00000
## Month OperatingSystems Browser Region
## Length:12199 Min. :1.000 Min. : 1.000 Min. :1.000
## Class :character 1st Qu.:2.000 1st Qu.: 2.000 1st Qu.:1.000
## Mode :character Median :2.000 Median : 2.000 Median :3.000
## Mean :2.124 Mean : 2.358 Mean :3.153
## 3rd Qu.:3.000 3rd Qu.: 2.000 3rd Qu.:4.000
## Max. :8.000 Max. :13.000 Max. :9.000
## TrafficType VisitorType Weekend Revenue
## Min. : 1.000 Length:12199 Mode :logical Mode :logical
## 1st Qu.: 2.000 Class :character FALSE:9343 FALSE:10291
## Median : 2.000 Mode :character TRUE :2856 TRUE :1908
## Mean : 4.075
## 3rd Qu.: 4.000
## Max. :20.000
```

*#getting measure of dispersion fro each cols*

*#Create a function*

```
library(moments)
summary.list = function(x)list(
  Mean=mean(x, na.rm=TRUE),
  Median=median(x, na.rm=TRUE),
  Skewness=skewness(x, na.rm=TRUE),
  Kurtosis=kurtosis(x, na.rm=TRUE),
  Variance=var(x, na.rm=TRUE),
  Std.Dev=sd(x, na.rm=TRUE),
  Coeff.Variation.Prcnt=sd(x, na.rm=TRUE)/mean(x, na.rm=TRUE)*100,
  Std.Error=sd(x, na.rm=TRUE)/sqrt(length(x[!is.na(x)]))
)
```

*#Calling the function and applying the function*

```
sapply(df_new[,c(y)], summary.list)
```

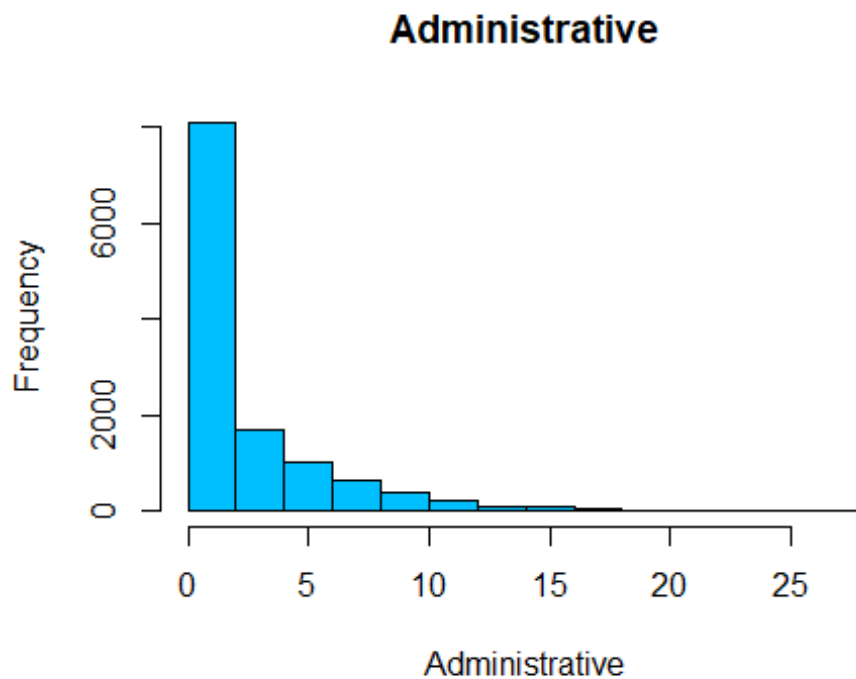
```
## Administrative Administrative_Duration Informational
## Mean 2.340028 81.68214 0.5088122
## Median 1 9 0
## Skewness 1.946248 5.59021 4.013451
## Kurtosis 7.636106 53.09389 29.64254
## Variance 11.09457 31516.25 1.62771
## Std.Dev 3.330851 177.5282 1.275817
## Coeff.Variation.Prcnt 142.3424 217.3402 250.7442
## Std.Error 0.03015735 1.60733 0.01155118
## Informational_Duration ProductRelated
## Mean 34.83734 32.05845
## Median 0 18
## Skewness 7.537435 4.332134
## Kurtosis 78.46409 34.04903
## Variance 20010.51 1989.241
## Std.Dev 141.4585 44.60091
```

```
## Coeff.Variation.Prcnt 406.0543          139.1237
## Std.Error            1.280758          0.4038142
##                      ProductRelated_Duration BounceRates ExitRates
## Mean                1207.508            0.02044674  0.04149678
## Median              609.5417            0.002930403  0.025
## Skewness            7.251403            3.152874    2.233125
## Kurtosis            139.5908            12.25506    7.624252
## Variance            3686121            0.002061387  0.0021388
## Std.Dev            1919.927            0.0454025    0.04624716
## Coeff.Variation.Prcnt 158.9991          222.0526    111.4476
## Std.Error          17.38292            0.0004110718 0.0004187193
##                      PageValues SpecialDay OperatingSystems Browser
## Mean                5.9525            0.06197229 2.124354    2.358144
## Median              0                0            2            2
## Skewness            6.348663            3.284481    2.031955    3.215653
## Kurtosis            67.94031            12.78605    13.26887    15.53659
## Variance            348.1132            0.03988432 0.8226229    2.926075
## Std.Dev            18.65779            0.1997106    0.9069856    1.710578
## Coeff.Variation.Prcnt 313.4446            322.2579    42.69465    72.53914
## Std.Error          0.1689266            0.001808169 0.008211799    0.01548748
##                      Region TrafficType
## Mean                3.153291            4.074596
## Median              3                2
## Skewness            0.9787304            1.958522
## Kurtosis            2.840195            6.466127
## Variance            5.771712            16.12675
## Std.Dev            2.402439            4.015813
## Coeff.Variation.Prcnt 76.18829            98.55732
## Std.Error          0.02175155 0.03635895
```

#Plots

```
library(tidyverse)
# Histograms for Area Income
hist(df_new$Administrative,
     main = "Administrative",
     xlab = "Administrative",
     col = "deepskyblue")
```





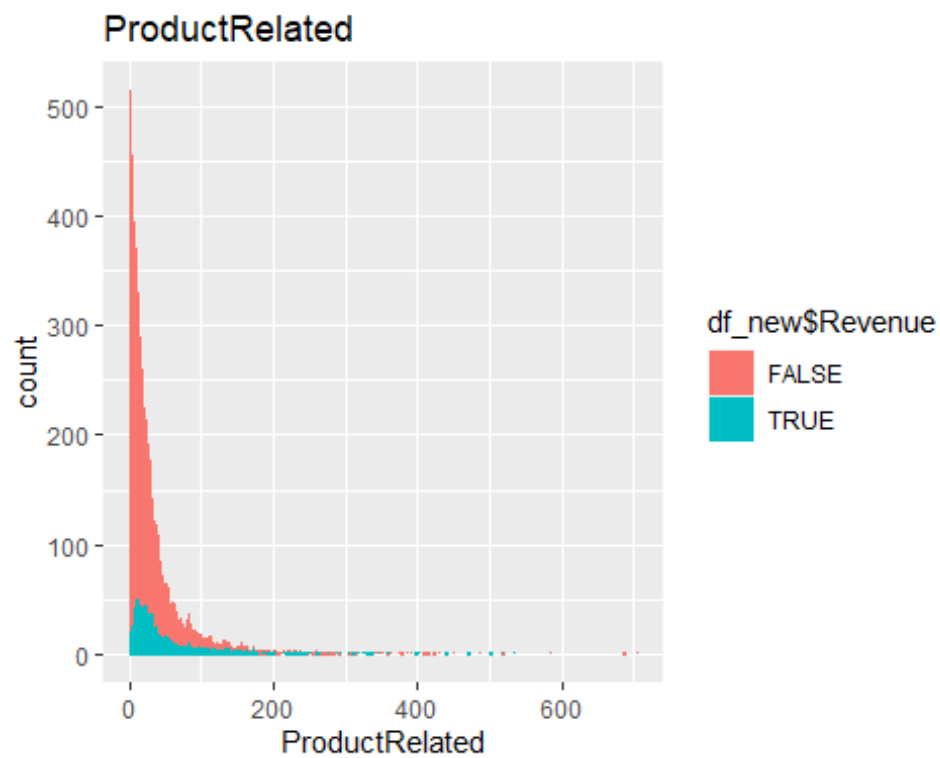
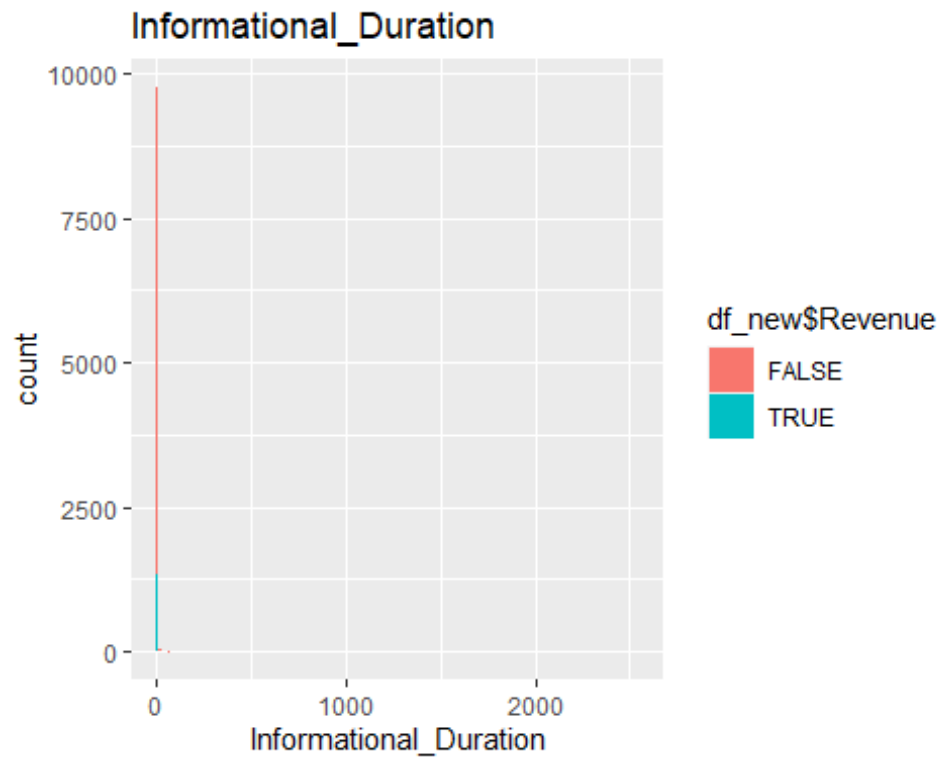



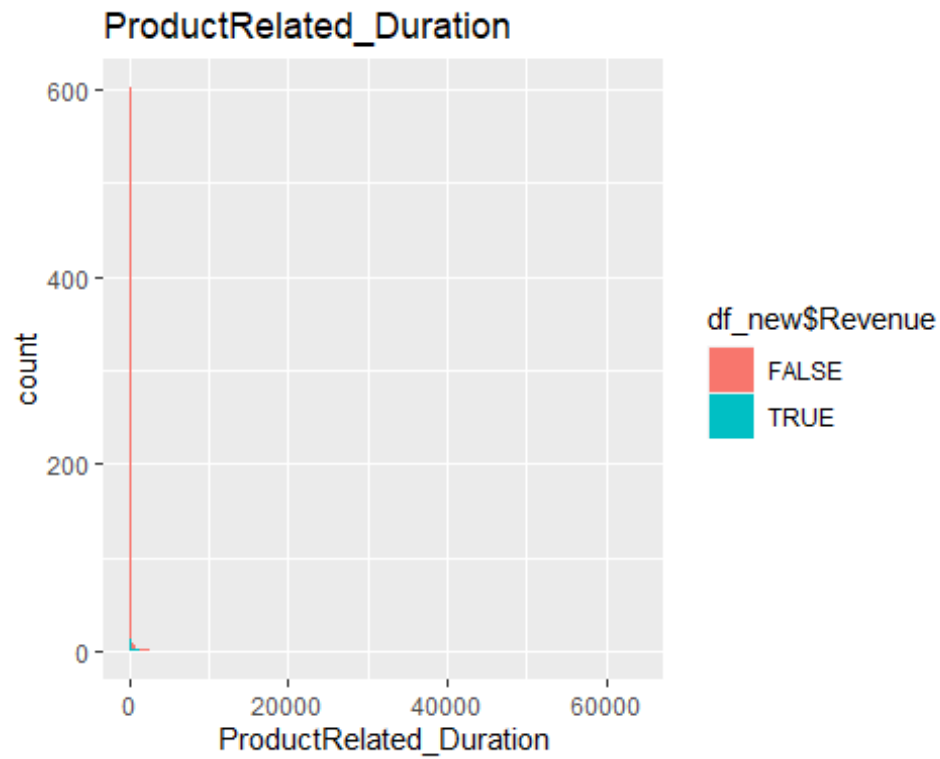
## Bivariant


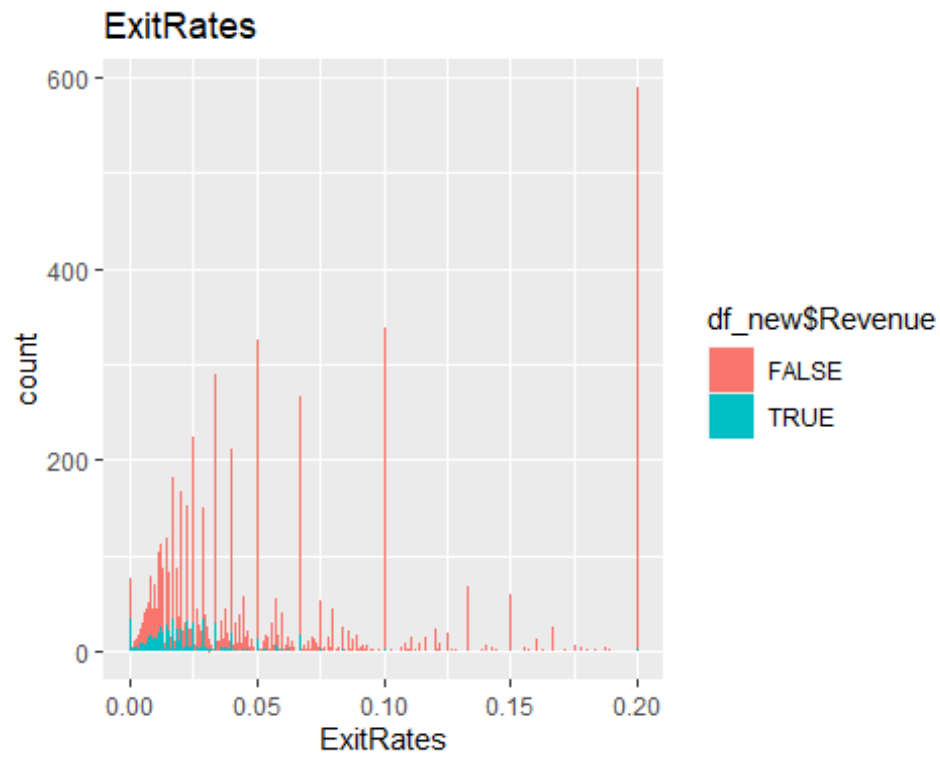
Analysis


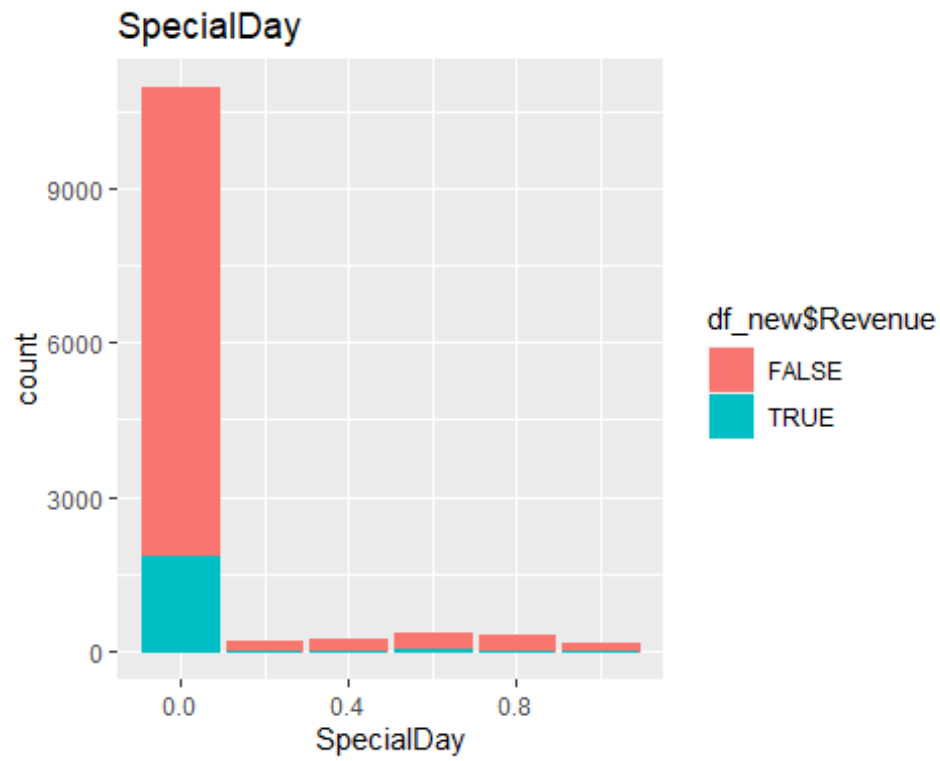
```
#We shall use loops to visuize how each column behave aganist revenue
for(i in 2:ncol(num)) { # Printing ggplot within
  for-loop
  print(ggplot(num, aes(x= num[ , i], fill = df_new$Revenue, color =
df_new$Revenue, )) +
    geom_bar()+labs(title =
'df_new[i]')+labs(title=colnames(df_new[i]), x=colnames(df_new[i])))
}
```


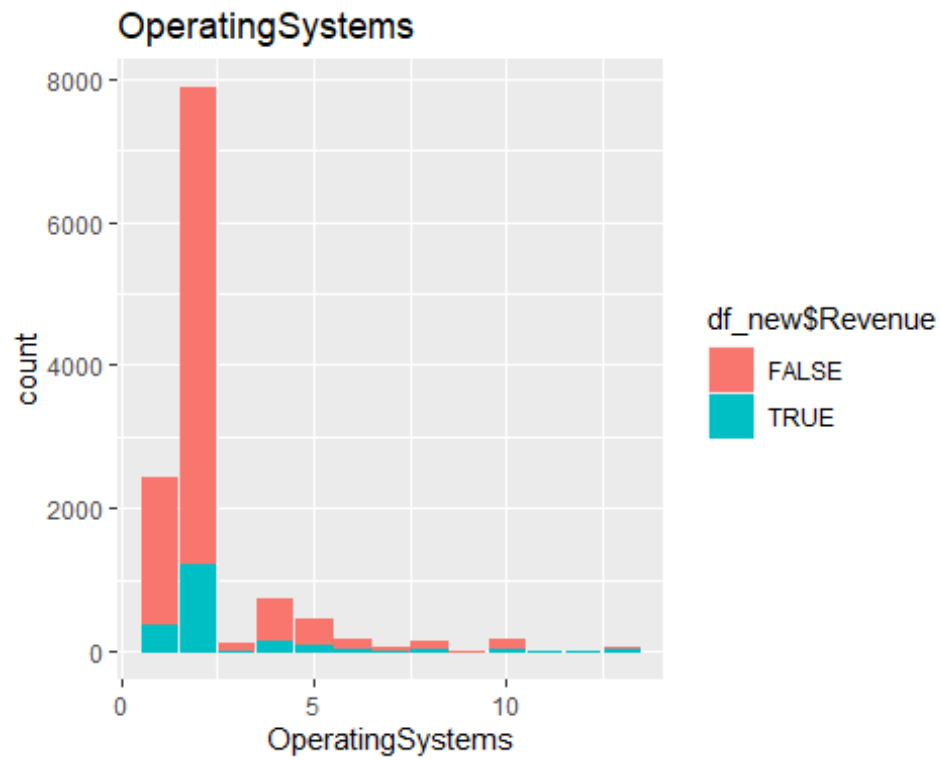


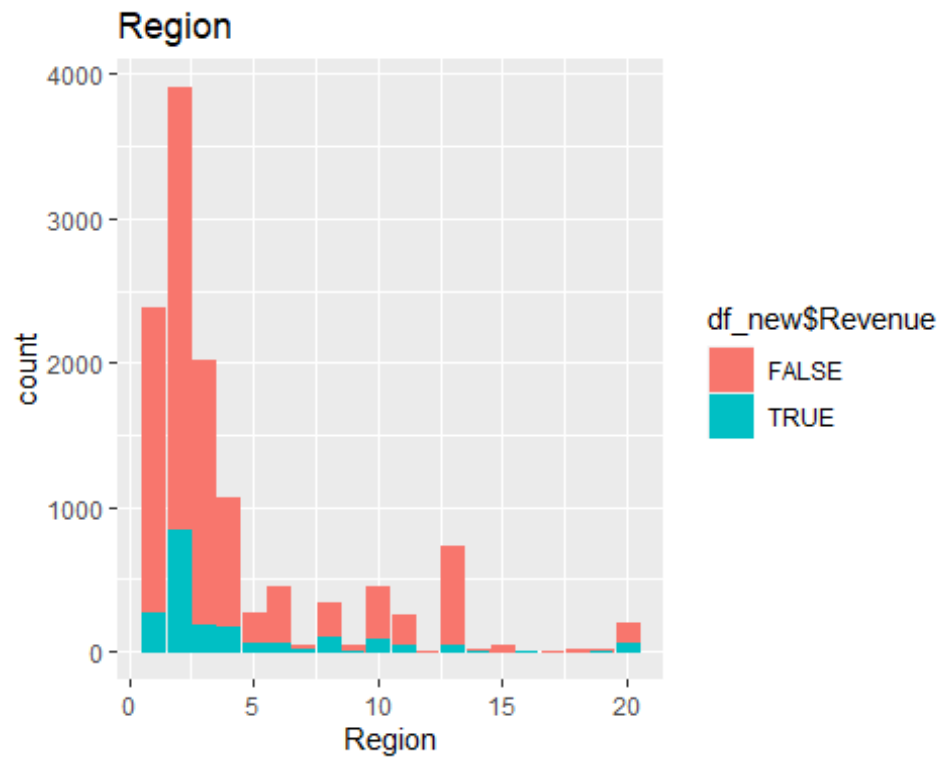













### Categorical months

*# Visualize revenue against months*


```
barplot(table(df_new$Revenue, df_new$Month), main = "Revenue per Month", col  
= c("orange", "green"), beside = TRUE,  
legend = rownames(table(df_new$Revenue, df_new$Month)), ylab="revenue", xlab  
= "Month")
```






November returns the highest number of revenues while February returns the lowest.

```
# Visualize revenue against Operating System
barplot(table(df_new$Revenue, df_new$OperatingSystems),
        main = "Revenue per Operating System",
        col = c('Orange', "green"), beside = TRUE,
        legend = rownames(table(df_new$Revenue, df_new$OperatingSystems)),
        ylab="revenue",
        xlab = "Operating System")
```




Operating System 2 returns the highest number of revenue while OS 5, 6, and 7 return the lowest.

```
# plotting the distribution of Revenue per Browser
barplot(table(df_new$Revenue, df_new$Browser),
        main = "Revenue per Browser",
        col = c("orange", "green"),
        beside = TRUE,
        legend = rownames(table(df_new$Revenue, df_new$Browser)),
        ylab="revenue",
        xlab = "Browser")
```




Browser 2 returns the highest number of revenue while 3, 7, 9, 11, and 12 return the lowest.

```
# plotting the distribution of Revenue per Region
barplot(table(df_new$Revenue, df_new$Region),
        main = "Revenue per Region",
        col = c("orange", "green"), beside = TRUE,
        legend = rownames(table(df_new$Revenue, df_new$Region)),
        ylab="revenue", xlab = "Region")
```




Region 1 returns the highest number of revenue, Region 5 and 8 returns the lowest.

```
# plotting the distribution of Revenue per Traffic Type
barplot(table(df_new$Revenue, df_new$TrafficType),
        main = "Revenue per Traffic Type",
        col = c("orange", "green"), beside = TRUE,
        legend = rownames(table(df_new$Revenue, df_new$TrafficType)),
        ylab="revenue", xlab = "Traffic Type")
```



Traffic 2 has the highest number of revenues, 12, 14 and 18 return the lowest.


```
# plotting the distribution of Revenue per Visitor Type
barplot(table(df_new$Revenue, df_new$VisitorType),
        main = "Revenue per Visitor Type",
        col = c("orange", "green"), beside = TRUE,
        legend = rownames(table(df_new$Revenue, df_new$VisitorType)),
        ylab="revenue", xlab = "Visitor Type")
```



Returning visitors

brought more revenue with new visitors generating around 1000.

```
# plotting the distribution of Revenue per Weekend
barplot(table(df_new$Revenue, df_new$Weekend),
        main = "Revenue per Weekend",
        col = c("orange", "green"), beside = TRUE,
        legend = rownames(table(df_new$Revenue, df_new$Weekend)),
        ylab="revenue", xlab = "Weekend")
```



More revenue was generated during the weekdays than the weekends.

*#check the correlation*

```
cor(df_new[,unlist(lapply(df_new, is.numeric))])
```

##	Administrative	Administrative_Duration	
Informational			
## Administrative	1.000000000	0.600409653	
0.37528761			
## Administrative_Duration	0.600409653	1.000000000	
0.30143630			
## Informational	0.375287611	0.301436296	
1.00000000			
## Informational_Duration	0.254786021	0.237189860	
0.61867795			
## ProductRelated	0.428191515	0.286783914	
0.37260472			
## ProductRelated_Duration	0.371027224	0.353513793	
0.38608372			
## BounceRates	-0.213666635	-0.137333397	-
0.10950530			
## ExitRates	-0.311274132	-0.202024452	-
0.15956681			
## PageValues	0.096920968	0.066168365	
0.04739015			
## SpecialDay	-0.097072098	-0.074736885	-
0.04937677			

## OperatingSystems 0.00962587	-0.006697922	-0.007610715	-
## Browser 0.03876681	-0.025763658	-0.015833675	-
## Region 0.03047732	-0.007262053	-0.006723711	-
## TrafficType 0.03518669	-0.034784126	-0.015075015	-
##	Informational_Duration	ProductRelated	
## Administrative	0.254786021	0.428191515	
## Administrative_Duration	0.237189860	0.286783914	
## Informational	0.618677947	0.372604721	
## Informational_Duration	1.000000000	0.279061948	
## ProductRelated	0.279061948	1.000000000	
## ProductRelated_Duration	0.346580691	0.860308186	
## BounceRates	-0.070159472	-0.193515772	
## ExitRates	-0.102932678	-0.286163211	
## PageValues	0.030064160	0.054115494	
## SpecialDay	-0.031293040	-0.025930622	
## OperatingSystems	-0.009749983	0.004090351	
## Browser	-0.019609349	-0.013706213	
## Region	-0.027920098	-0.040106501	
## TrafficType	-0.025163571	-0.044344333	
##	ProductRelated_Duration	BounceRates	ExitRates
## Administrative	0.371027224	-0.213666635	-0.311274132
## Administrative_Duration	0.353513793	-0.137333397	-0.202024452
## Informational	0.386083717	-0.109505298	-0.159566815
## Informational_Duration	0.346580691	-0.070159472	-0.102932678
## ProductRelated	0.860308186	-0.193515772	-0.286163211
## ProductRelated_Duration	1.000000000	-0.174375499	-0.245334012
## BounceRates	-0.174375499	1.000000000	0.903358192
## ExitRates	-0.245334012	0.903358192	1.000000000
## PageValues	0.050840624	-0.115991977	-0.173571542
## SpecialDay	-0.038210652	0.087839995	0.116783762
## OperatingSystems	0.002775788	0.026839839	0.016482012
## Browser	-0.007838332	-0.016018380	-0.003565541
## Region	-0.034862498	0.001432015	-0.001837556
## TrafficType	-0.037506944	0.089199039	0.087386232
##	PageValues	SpecialDay	OperatingSystems
Browser			
## Administrative 0.025763658	0.09692097	-0.097072098	-0.006697922 -
## Administrative_Duration 0.015833675	0.06616837	-0.074736885	-0.007610715 -
## Informational 0.038766808	0.04739015	-0.049376774	-0.009625870 -
## Informational_Duration 0.019609349	0.03006416	-0.031293040	-0.009749983 -
## ProductRelated 0.013706213	0.05411549	-0.025930622	0.004090351 -



```
## ProductRelated_Duration 0.05084062 -0.038210652 0.002775788 -
0.007838332
## BounceRates -0.11599198 0.087839995 0.026839839 -
0.016018380
## ExitRates -0.17357154 0.116783762 0.016482012 -
0.003565541
## PageValues 1.00000000 -0.064532709 0.018583782
0.045845065
## SpecialDay -0.06453271 1.000000000 0.012757766
0.003465984
## OperatingSystems 0.01858378 0.012757766 1.000000000
0.212244823
## Browser 0.04584506 0.003465984 0.212244823
1.000000000
## Region 0.01059087 -0.016452464 0.071953240
0.091889464
## TrafficType 0.01223694 0.052827944 0.182874100
0.102886237
## Region TrafficType
## Administrative -0.007262053 -0.03478413
## Administrative_Duration -0.006723711 -0.01507502
## Informational -0.030477323 -0.03518669
## Informational_Duration -0.027920098 -0.02516357
## ProductRelated -0.040106501 -0.04434433
## ProductRelated_Duration -0.034862498 -0.03750694
## BounceRates 0.001432015 0.08919904
## ExitRates -0.001837556 0.08738623
## PageValues 0.010590868 0.01223694
## SpecialDay -0.016452464 0.05282794
## OperatingSystems 0.071953240 0.18287410
## Browser 0.091889464 0.10288624
## Region 1.000000000 0.04252523
## TrafficType 0.042525234 1.00000000
```


```
#install.packages("corrplot")
library(corrplot)
```

```
## corrplot 0.90 loaded
```

```
#
```

```
## Let's build a correlation matrix to understand the relation between each attributes
```

```
corrplot(cor(num), method = 'number')
```



```
#drop cols highly correlated
col_drop <- c("Administrative_Duration", "Informational_Duration",
"ProductRelated_Duration", "ExitRates")
df_new <- df_new[, !names(df_new) %in% col_drop]
head(df_new)

##   Administrative Informational ProductRelated BounceRates PageValues
SpecialDay
## 1           0           0           1  0.20000000           0
0
## 2           0           0           2  0.00000000           0
0
## 3           0           0           1  0.20000000           0
0
## 4           0           0           2  0.05000000           0
0
## 5           0           0          10  0.02000000           0
0
## 6           0           0          19  0.01578947           0
0
##   Month OperatingSystems Browser Region TrafficType   VisitorType
Weekend
## 1   Feb               1       1       1           1 Returning_Visitor
FALSE
## 2   Feb               2       2       1           2 Returning_Visitor
FALSE
## 3   Feb               4       1       9           3 Returning_Visitor
```

```

FALSE
## 4 Feb 3 2 2 4 Returning_Visitor
FALSE
## 5 Feb 3 3 1 4 Returning_Visitor
TRUE
## 6 Feb 2 2 1 3 Returning_Visitor
FALSE
## Revenue
## 1 FALSE
## 2 FALSE
## 3 FALSE
## 4 FALSE
## 5 FALSE
## 6 FALSE

```

## Modelling

*#Check head*  
head(df\_new)

```

## Administrative Informational ProductRelated BounceRates PageValues
SpecialDay
## 1 0 0 1 0.20000000 0
0
## 2 0 0 2 0.00000000 0
0
## 3 0 0 1 0.20000000 0
0
## 4 0 0 2 0.05000000 0
0
## 5 0 0 10 0.02000000 0
0
## 6 0 0 19 0.01578947 0
0
## Month OperatingSystems Browser Region TrafficType VisitorType
Weekend
## 1 Feb 1 1 1 1 Returning_Visitor
FALSE
## 2 Feb 2 2 1 2 Returning_Visitor
FALSE
## 3 Feb 4 1 9 3 Returning_Visitor
FALSE
## 4 Feb 3 2 2 4 Returning_Visitor
FALSE
## 5 Feb 3 3 1 4 Returning_Visitor
TRUE
## 6 Feb 2 2 1 3 Returning_Visitor
FALSE
## Revenue
## 1 FALSE

```

```
## 2 FALSE
## 3 FALSE
## 4 FALSE
## 5 FALSE
## 6 FALSE
```

*#selecting data without revenue*

```
data<-df_new[,-14]
head(data)
```

```
## Administrative Informational ProductRelated BounceRates PageValues
SpecialDay
```

```
## 1      0      0      1 0.20000000      0
0
## 2      0      0      2 0.00000000      0
0
## 3      0      0      1 0.20000000      0
0
## 4      0      0      2 0.05000000      0
0
## 5      0      0     10 0.02000000      0
0
## 6      0      0     19 0.01578947      0
0
```

```
## Month OperatingSystems Browser Region TrafficType VisitorType
Weekend
```

```
## 1 Feb      1      1      1      1 Returning_Visitor
FALSE
## 2 Feb      2      2      1      2 Returning_Visitor
FALSE
## 3 Feb      4      1      9      3 Returning_Visitor
FALSE
## 4 Feb      3      2      2      4 Returning_Visitor
FALSE
## 5 Feb      3      3      1      4 Returning_Visitor
TRUE
## 6 Feb      2      2      1      3 Returning_Visitor
FALSE
```

*# Create custom function to fix data types and round*

```
to_numeric_and_round_func <- function(x){
  round(as.numeric(as.character(x)),2)
}
```

*# Mutate the columns to proper data type*

```
data <- data %>%
  mutate_at(vars(-one_of("Month", "Region", "VisitorType", "Weekend")),
to_numeric_and_round_func)
```

*# create clean data with no NA*

```
clean_data <- data %>%
  drop_na()
```

```
##Kmeans
```

```
# Set seed
```

```
set.seed(1234)
```

```
col.names<-c("Month", "VisitorType", "Weekend")
```

```
# Cluster Analysis - kmeans
```

```
kmeans_basic <- kmeans(clean_data[, !names(data) %in% col.names], centers = 5)
```

```
kmeans_basic_table <- data.frame(kmeans_basic$size, kmeans_basic$centers)
```

```
kmeans_basic_df <- data.frame(Cluster = kmeans_basic$cluster, data)
```

```
# head of df
```

```
head(kmeans_basic_df)
```

```
## Cluster Administrative Informational ProductRelated BounceRates  
PageValues
```

```
## 1      4      0      0      1      0.20  
0
```

```
## 2      4      0      0      2      0.00  
0
```

```
## 3      4      0      0      1      0.20  
0
```

```
## 4      4      0      0      2      0.05  
0
```

```
## 5      4      0      0     10      0.02  
0
```

```
## 6      4      0      0     19      0.02  
0
```

```
## SpecialDay Month OperatingSystems Browser Region TrafficType
```

```
## 1      0 Feb      1      1      1      1
```

```
## 2      0 Feb      2      2      1      2
```

```
## 3      0 Feb      4      1      9      3
```

```
## 4      0 Feb      3      2      2      4
```

```
## 5      0 Feb      3      3      1      4
```

```
## 6      0 Feb      2      2      1      3
```

```
## VisitorType Weekend
```

```
## 1 Returning_Visitor FALSE
```

```
## 2 Returning_Visitor FALSE
```

```
## 3 Returning_Visitor FALSE
```

```
## 4 Returning_Visitor FALSE
```

```
## 5 Returning_Visitor TRUE
```

```
## 6 Returning_Visitor FALSE
```


```
# Visulize the clusters per month
```

```
ggplot(data = kmeans_basic_df, aes(y = Cluster)) +
```

```
  geom_bar(aes(fill = Month)) +
```

```
  ggtitle("Count of Clusters by Month") +
```

```
  theme(plot.title = element_text(hjust = 0.5))
```



#### elbow method

```
library(factoextra)
```


```
## Warning: package 'factoextra' was built under R version 4.0.5
```

```
## Welcome! Want to learn more? See two factoextra-related books at  
https://goo.gl/ve3WBa
```

```
data_norm <- scale(clean_data[, !names(data) %in% col.names])
```

```
# Get the optimum number of clusters
```


```
fviz_nbclust(data_norm, kmeans, method = "wss")
```



The optimum clusters from the above is 3.

```
#kmeans
kmeans_fancy <- kmeans(data_norm, 3, nstart = 20)

# plot the clusters
fviz_cluster(kmeans_fancy, data = data_norm, geom = c("point"), ellipse.type =
"euclid")
```



*#Check the size of each cluster*

```
kmeans_fancy $size
```

```
## [1] 8885 1561 1753
```

The first cluster has 8885 values, second has 1561 and third has 1753 values

*# check their response to revenue*

```
table(kmeans_fancy$cluster, df_new$Revenue)
```

```
##
```

```
## FALSE TRUE
```

```
## 1 7519 1366
```

```
## 2 1506 55
```

```
## 3 1266 487
```

```
data %>%
```

```
  mutate(Cluster = kmeans_fancy$cluster) %>%
```

```
  group_by(Cluster) %>%
```

```
  summarize_all('median')
```

```
## # A tibble: 3 x 14
```

```
## Cluster Administrative Informational ProductRelated BounceRates  
PageValues
```

```
## <int> <dbl> <dbl> <dbl> <dbl>
```

```
<dbl>
```

```
## 1 1 1 0 16 0 0
```

```
## 2 2 0 0 5 0.07 0
```



```
## 3      3      7      2      70      0
2.09
## # ... with 8 more variables: SpecialDay <dbl>, Month <chr>,
## #   OperatingSystems <dbl>, Browser <dbl>, Region <int>, TrafficType
## #   <dbl>,
## #   VisitorType <chr>, Weekend <lgl>
```

## Hierarchical clustering

```
library(cluster)


## Warning: package 'cluster' was built under R version 4.0.5

# compute the euclidean distance using euclidean metric
eucl_dist<- dist(data_norm, method = "euclidean")
#compute hierarchical clustering using the Ward method
res_hc<- hclust(eucl_dist, method = "ward.D2")
res_hc

##
## Call:
## hclust(d = eucl_dist, method = "ward.D2")
##
## Cluster method      : ward.D2
## Distance            : euclidean
## Number of objects: 12199

# plot the obtained dendrogram
plot(res_hc, cex = 0.6, hang = -1)
```


## Cluster Dendrogram



```
eucl_dist  
hclust (*, "ward.D2")
```

```
# compute the euclidean distance using manhattan metric  
eucl_dist_man<- dist(data_norm, method = "manhattan")  
#compute hierarchical clustering using the Ward method  
res_hc_man<- hclust(eucl_dist_man, method = "ward.D2")  
res_hc_man  
  
##  
## Call:  
## hclust(d = eucl_dist_man, method = "ward.D2")  
##  
## Cluster method      : ward.D2  
## Distance             : manhattan  
## Number of objects: 12199  
  
# plot the obtained dendrogram  
plot(res_hc_man, cex = 0.6, hang = -1)
```

## Cluster Dendrogram



eucl\_dist\_man  
hclust (\*, "ward.D2")